

Linear Regression

May 13, 2019

Linear Regression with one variable

```
[1]: import numpy as np
import matplotlib.pyplot as plt
from numpy.linalg import pinv
%matplotlib inline
```

Let us define our hypothesis, costFunction and implementation of gradient Descent in a vectorised way..

```
[2]: def hypothesis(Theta,X):
    return X@Theta    #Vectorised Implementation
```

```
[3]: def costFunction(Theta,X,Y):
    m,n = X.shape
    return ((0.5/m)*((hypothesis(Theta,X) - Y).T @ (hypothesis(Theta,X) - Y))).
    →item()    #Vectorised Implementation
```

```
[4]: def gradientDescent(Theta,X,Y,alpha,iterations):
    m,n = X.shape
    for count in range(0,iterations):
        Theta = Theta - (alpha * (1/m)*(X.T @ (hypothesis(Theta,X)-Y)))
    →#Vectorised Implementation
    return Theta
```


We have the data in csv format as population,profit

```
[5]: data = np.loadtxt("Data/ex1data1.txt", delimiter=",")
X = data[:, :-1]
Y = data[:, -1:]
```

Visualizing data: Population vs Profit

```
[6]: fig, ax = plt.subplots()
ax.plot(X,Y,ls='', marker = 'x', label='Population vs Profit')
ax.set(xlabel="Population of city in 10,000's", ylabel="Profits from the city_
    →in $10,000", title='Population vs Profit')
```

```
[6]: [Text(0, 0.5, 'Profits from the city in $10,000'),
Text(0.5, 0, "Population of city in 10,000's"),
Text(0.5, 1.0, 'Population vs Profit')]
```



Regression_files/LinearRegression91.bb

Converting data into matrices as adding column of ones to X

```
[7]: X = np.matrix(X)
Y = np.matrix(Y)
m,n = X.shape
X_new = np.hstack((np.ones((m,1)),X))
n = n+1
```

Let us take, theta initially as zeros and set learning rate (alpha) and number of iterations

```
[8]: initialTheta = np.zeros((n,1))
alpha = 0.01 #Learning rate
iterations = 1500
```

Cost when initial parameters are zero.

```
[9]: print("Cost when theta is\n",initialTheta,"\nis:␣
→",costFunction(initialTheta,X_new,Y))
```

Cost when theta is

```
[[0.]
 [0.]]
is: 32.072733877455676
```

Running Gradient Descent:

```
[10]: finalTheta = gradientDescent(initialTheta,X_new,Y,alpha,iterations)
print("Optimized parameters after running gradient Descent are: \n",finalTheta)
```

Optimized parameters after running gradient Descent are:

```
[-3.63029144]
[ 1.16636235]]
```

Cost after running Gradient Descent:


```
[11]: print("Cost after running gradient Descent is:␣
      →",costFunction(finalTheta,X_new,Y))
```

Cost after running gradient Descent is: 4.483388256587725

Linear fit after obtaining the values of theta:

```
[12]: fig, ax = plt.subplots()
      ax.plot(X,Y,ls='', marker = 'x', label='Population vs Profit')
      ax.plot(X,hypothesis(finalTheta,X_new))
      ax.set(xlabel="Population of city in 10,000's", ylabel="Profits from the city␣
      →in $10,000", title='Population vs Profit')
```

```
[12]: [Text(0, 0.5, 'Profits from the city in $10,000'),
      Text(0.5, 0, "Population of city in 10,000's"),
      Text(0.5, 1.0, 'Population vs Profit')]
```



Regression_files/LinearRegression21.bb

Predicting profits for given population using our model:

```
[13]: print("Profits in $10,000's when the city population is 35,000: ",[1, 3.
      →5]*finalTheta)
      print("Profits in $10,000's when the city population is 70,000: ",[1,␣
      →7]*finalTheta)
```

Profits in \$10,000's when the city population is 35,000: [[0.45197679]]

Profits in \$10,000's when the city population is 70,000: [[4.53424501]]

Linear Regression with multivariable

Note: hypothesis, costFunction, gradientDescent are same as above.

Feature Normalization:

```
[14]: def featureNormalization(X):
    m, n = X.shape
    means = []
    stds = []
    for column in range(1,n):
        mean = X[:,column].mean()
        std = X[:,column].std()
        means.append(mean)
        stds.append(std)
        X[:,column] = (X[:,column] - mean)/std
    return X, means, stds
```

```
[15]: def predict(X,theta,means,stds):
    X = np.matrix(X)
    m,n = X.shape
    for column in range(1,n):
        X[0:,column] = (X[0:,column] - means[column-1])/stds[column-1]
    return X@theta
```

Loading Data:

```
[16]: data = np.loadtxt("Data/ex1data2.txt", delimiter=',')
X = data[:, :-1]
Y = data[:, -1:]
```

Converting data into matrices, adding column of ones to X and normalizing features:

```
[17]: X = np.matrix(X)
Y = np.matrix(Y)
m,n = X.shape
X_new = np.hstack((np.ones((m,1)),X))
X_new,means,stds = featureNormalization(X_new)
n = n+1
```

Intializing theta, learning rate and number of iterations:

```
[18]: initialTheta = np.zeros((n,1))
alpha = 0.01 #Learning rate
iterations = 400
```

Cost when initial parameters are zero.

```
[19]: print("Cost when theta is\n",initialTheta,"\nis:␣
→",costFunction(initialTheta,X_new,Y))
```

Cost when theta is

```
[[0.]
 [0.]
 [0.]]
is: 65591548106.45744
```

Running Gradient Descent:

```
[20]: finalTheta = gradientDescent(initialTheta,X_new,Y,alpha,iterations)
print("Optimized parameters after running gradient Descent are: \n",finalTheta)
```

Optimized parameters after running gradient Descent are:

```
[[334302.06399328]
 [ 99411.44947359]
 [ 3267.01285407]]
```

Cost after running Gradient Descent:

```
[21]: print("Cost after running gradient Descent is:␣
      ↪",costFunction(finalTheta,X_new,Y))
```

Cost after running gradient Descent is: 2105448288.6292472

Predicting cost of the houses using our model:

```
[22]: print("Predicted price of a 1650 sq-ft, 3 br house (using gradient descent):␣
      ↪",predict([1.,1650.,3.],finalTheta,means,stds).item())
```

Predicted price of a 1650 sq-ft, 3 br house (using gradient descent):
289221.5473712181

Normal equation Method

```
[23]: def normalEqn(X,y):
      ↪ return pinv(X.T@X)@X.T@y
```

Loading Data

```
[24]: data = np.loadtxt("Data/ex1data2.txt", delimiter=',')
X = data[:, :-1]
Y = data[:, -1:]
```

```
[25]: X = np.matrix(X)
Y = np.matrix(Y)
m,n = X.shape
X_new = np.hstack((np.ones((m,1)),X))
n = n+1
```

Applying normal equation to get theta

```
[26]: theta = normalEqn(X_new,Y)
```

Predicting values:

```
[27]: print("Predicted price of a 1650 sq-ft, 3 br house (using normal eqn): ",[1.
      ↪,1650.,3.]*theta)
```

Predicted price of a 1650 sq-ft, 3 br house (using normal eqn):
[[293081.46433497]]